

CSCE 735 Fall 2022

Name: Rohan Chaudhury

UIN: 432001358

HW 2: Parallel Merge Sort Using Threads

1. (70 points) Revise the code to implement a thread-based parallel merge sort. The code should compile successfully and should report error=0 for the following instances:

./sort_list.exe 4 1

./sort_list.exe 4 2

./sort_list.exe 4 3

./sort_list.exe 20 4

./sort_list.exe 24 8

ANSWER:

The parallelized code is attached in the zip file. The logs are as follows:

List Size = 16, Threads = 2, error = 0, time (sec) = 0.0005, qsort_time = 0.0000

List Size = 16, Threads = 4, error = 0, time (sec) = 0.0006, qsort_time = 0.0000

List Size = 16, Threads = 8, error = 0, time (sec) = 0.0009, qsort_time = 0.0000

List Size = 1048576, Threads = 16, error = 0, time (sec) = 0.0217, qsort_time = 0.1773


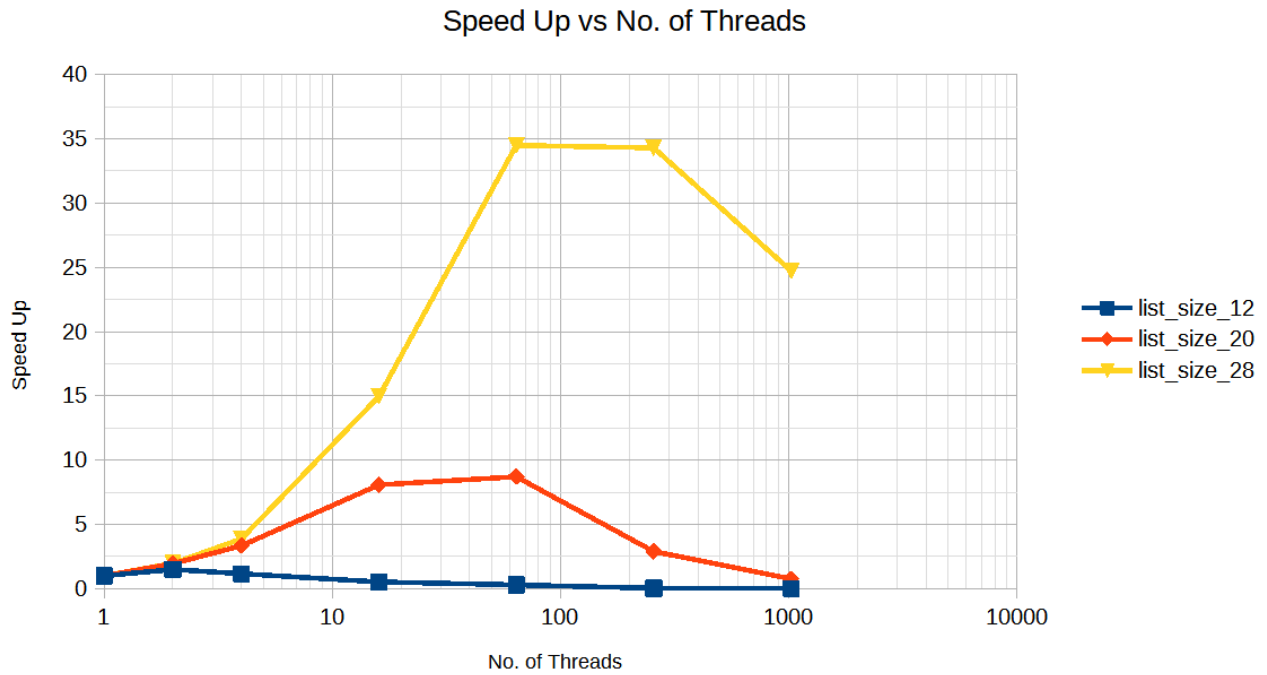
List Size = 16777216, Threads = 256, error = 0, time (sec) = 0.2168, qsort_time = 3.3305

| | | | | | | | | | |
|-----------|----------|---------|-----|-------|---|------------|--------|------------|--------|
| List Size | 16 | Threads | 2 | error | 0 | time (sec) | 0.0005 | qsort_time | 0 |
| List Size | 16 | Threads | 4 | error | 0 | time (sec) | 0.0006 | qsort_time | 0 |
| List Size | 16 | Threads | 8 | error | 0 | time (sec) | 0.0009 | qsort_time | 0 |
| List Size | 1048576 | Threads | 16 | error | 0 | time (sec) | 0.0217 | qsort_time | 0.1773 |
| List Size | 16777216 | Threads | 256 | error | 0 | time (sec) | 0.2168 | qsort_time | 3.3305 |

As can be observed, the error returned for all of the above tests is 0. It should also be emphasized that for smaller list sizes, parallelization has no advantage over single threaded processing. For list sizes 20 and 24, the predicted performance increase over single threaded processing is seen.

2. (20 points) Plot speedup and efficiency for all combinations of k and q chosen from the following sets: k = 12, 20, 28 ; q = 0, 1, 2, 4, 6, 8, 10. Comment on how the results of your experiments align with or diverge from your understanding of the expected behavior of the parallelized code.

ANSWER:



We can see the following things from the two graphs above:

1. When $k=12$ (that is, list size = 2^{12}) is compared to $k=20$ or $k=28$, the speedup is quite small. Because employing many threads adds thread management cost when the input data size is comparatively small. When the time required to run the thread is similar to this cost, the efficiency of serial execution decreases.
2. The speedup increases till a point and then decreases for $k=20$ and $k=28$. This is due to the excessive number of threads, and the thread management cost begins to play a significant part in time consumed.

- The number of threads required to achieve maximum speed is varied between $k=20$ and $k=28$. This is due to the work being divided into smaller jobs for the threads to do, resulting in a shorter overall time spent. However, after reaching the maximum speedup, raising the number of threads increases the time taken due to context switching across the threads.
- The necessity for thread synchronization and resource management reduces efficiency as the number of threads increases.

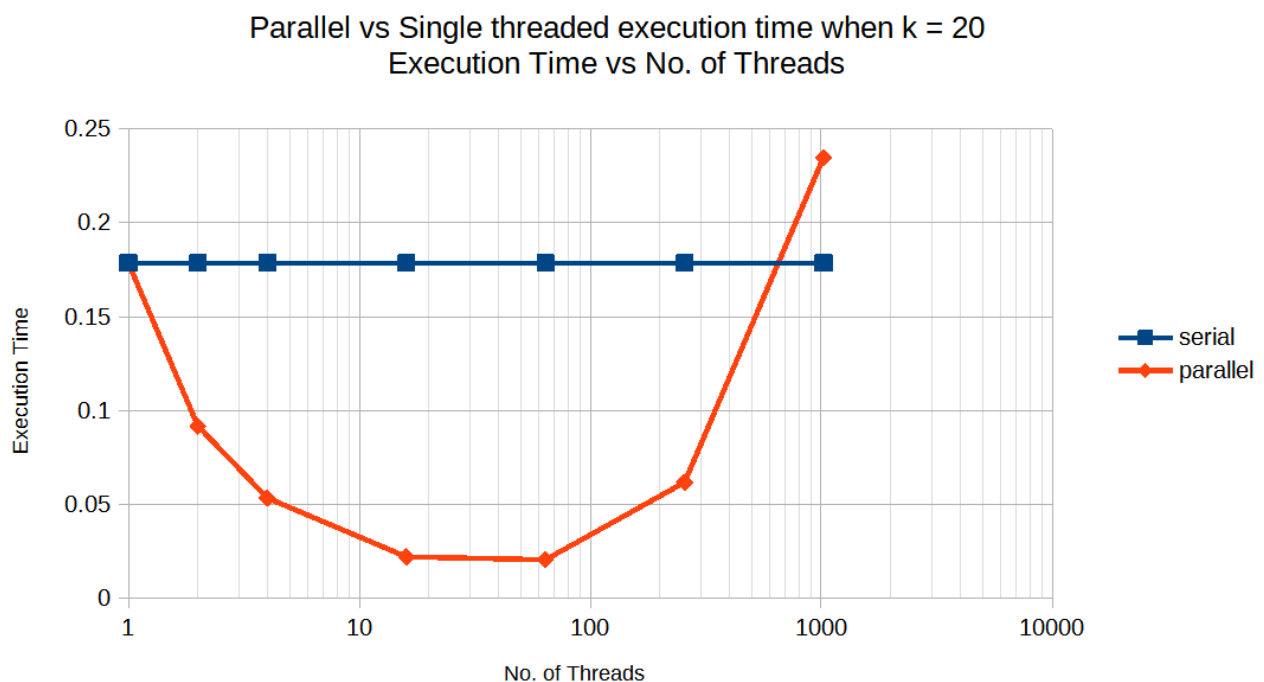
3. (10 points) Your code should demonstrate speedup when sorting lists of appropriate sizes. Determine two values of k for which your code shows speedup as q is varied. Present the timing results for your code along with speedup and efficiency obtained to convince the reader that you have a well-designed parallel merge sort. You may use results from experiments in previous problems or identify new values k and q to illustrate how well your code has been parallelized.

ANSWER:

As q was varied, two k values for speedup were observed: 20, 28.

Below is the table for $k=20$


| | | | | | | | | | | | | | |
|-----------|---------|---------|------|-------|---|------------|--------|------------|--------|------------|-------------|--------------|-------------|
| List Size | 1048576 | Threads | 1 | error | 0 | time (sec) | 0.1785 | qsort_time | 0.1743 | Speed up = | 1 | Efficiency = | 1 |
| List Size | 1048576 | Threads | 2 | error | 0 | time (sec) | 0.0917 | qsort_time | 0.1769 | Speed up = | 1.946564885 | Efficiency = | 0.973282443 |
| List Size | 1048576 | Threads | 4 | error | 0 | time (sec) | 0.0534 | qsort_time | 0.1744 | Speed up = | 3.342696629 | Efficiency = | 0.835674157 |
| List Size | 1048576 | Threads | 16 | error | 0 | time (sec) | 0.0221 | qsort_time | 0.1767 | Speed up = | 8.076923077 | Efficiency = | 0.504807692 |
| List Size | 1048576 | Threads | 64 | error | 0 | time (sec) | 0.0205 | qsort_time | 0.1743 | Speed up = | 8.707317073 | Efficiency = | 0.136051829 |
| List Size | 1048576 | Threads | 256 | error | 0 | time (sec) | 0.0617 | qsort_time | 0.1747 | Speed up = | 2.893030794 | Efficiency = | 0.011300902 |
| List Size | 1048576 | Threads | 1024 | error | 0 | time (sec) | 0.2348 | qsort_time | 0.1729 | Speed up = | 0.760221465 | Efficiency = | 0.000742404 |



As demonstrated in the table and chart above, with $k = 20$, execution time falls with parallelization as the number of threads increases until it reaches the best option (at number of threads = $2^6 = 64$) with the least runtime, 0.0205 seconds. Raising the number of threads has reduction effect on the execution time after this point. The graph compares execution durations in parallel mode (orange line) with single threaded mode (blue line) over q values, which represent the number of threads used.

Below is the table for $k = 28$

| | | | | | | | | | | | |
|-----------|-----------|---------|------------|--------------|---------|------------|---------|------------|-------------|--------------|-------------|
| List Size | 268435456 | Threads | 1 error | 0 time (sec) | 62.9848 | qsort_time | 62.5606 | Speed up = | 1 | Efficiency = | 1 |
| List Size | 268435456 | Threads | 2 error | 0 time (sec) | 31.9119 | qsort_time | 62.5535 | Speed up = | 1.973708867 | Efficiency = | 0.986854434 |
| List Size | 268435456 | Threads | 4 error | 0 time (sec) | 16.2256 | qsort_time | 62.5607 | Speed up = | 3.881816389 | Efficiency = | 0.970454097 |
| List Size | 268435456 | Threads | 16 error | 0 time (sec) | 4.2106 | qsort_time | 62.5397 | Speed up = | 14.95862822 | Efficiency = | 0.934914264 |
| List Size | 268435456 | Threads | 64 error | 0 time (sec) | 1.8261 | qsort_time | 62.891 | Speed up = | 34.49142982 | Efficiency = | 0.538928591 |
| List Size | 268435456 | Threads | 256 error | 0 time (sec) | 1.8367 | qsort_time | 62.5081 | Speed up = | 34.29237219 | Efficiency = | 0.133954579 |
| List Size | 268435456 | Threads | 1024 error | 0 time (sec) | 2.5465 | qsort_time | 63.1247 | Speed up = | 24.73387002 | Efficiency = | 0.02415417 |



As depicted in the table and graph above, for $k = 28$, execution time falls with parallelization as the number of threads increases until it reaches the best choice (at number of threads = $2^6 = 64$) with the shortest runtime, i.e., 1.8261 sec. After this, increasing the number of threads has no effect on execution time reduction, as seen by the progressive increase in execution times beyond number of threads = $2^6 = 64$. The graph compares execution durations in parallel mode (orange line) with single threaded mode (blue line) over q values, which represent the number of threads used.